

COMP9417 HomeWork 2

Ping GAO z5163482

March 29, 2020

1 Q1

1.1 Part A

DecisionTreeClassifier										
Dataset	5%	10%	15%	20%	25%	30%	35%	40%	45%	50%
australian	72.61%	74.35%	75.36%	77.39%	77.83%	79.71%	83.77%	81.16%	80.72%	83.48%
balance-scale	69.92%	75.04%	69.12%	74.24%	74.40%	75.52%	78.08%	75.68%	77.92%	76.64%
hypothyroid	94.94%	96.31%	97.77%	99.18%	99.20%	99.42%	99.42%	99.52%	99.34%	99.20%

BernoulliNB with priors										
Dataset	5%	10%	15%	20%	25%	30%	35%	40%	45%	50%
australian	73.48%	79.86%	81.45%	80.43%	79.71%	79.86%	79.86%	81.16%	82.17%	81.88%
balance-scale	46.08%	46.08%	46.08%	46.08%	46.24%	46.08%	46.08%	46.24%	46.24%	46.08%
hypothyroid	91.38%	91.81%	92.23%	92.23%	92.23%	92.26%	92.23%	92.23%	92.23%	92.23%

1.2 Part B

I think (3), (5) statements are true.

1.3 Part C

I choose (1).

2 Q2

2.1 Part A

My accuracy score for the test dataset is 82.77%.

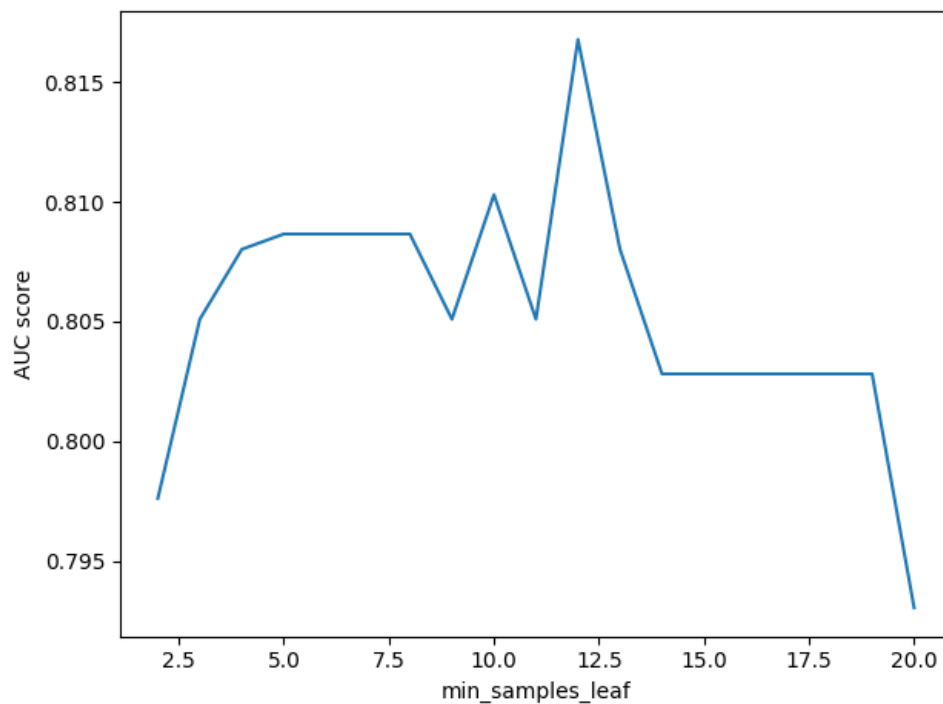
My accuracy score for the training dataset is 85.65%.

2.2 Part B

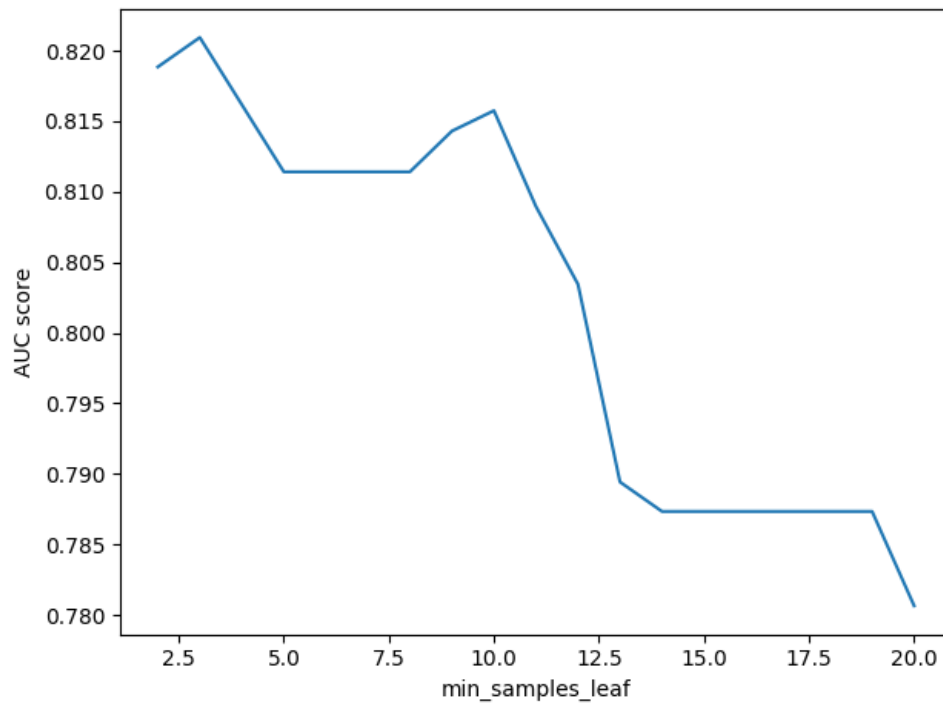
The min.samples_leaf number of 5 to 7 give the optimal result, this can be observed by compare the auc score in the part c's plots, pick the maximum score with the lowest variance.

2.3 Part C

Plot For Test Dataset



Plot For Train Dataset



2.4 Part D

We make the assumption that 'Sex' and 'Pclass' are independent feature.

Thus, $(S=true | G = female, C = 1) = P(S = true | G = female) * P(S = true | C = 1)$

$P(S = true | G = female) = P(G = female) \cap P(S = true) / P(G = female)$

$= 109/573$

$P(S = true | C = 1) = 136/216$

$P(S = true | G = female, C = 1) = 11.98\%$.

2.5 My code

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import tree
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score, roc_curve, auc

# import numpy as np
# import seaborn as sns
```

```

# import graphviz
# import matplotlib.pyplot as plt

class TitanicSinkingModel:
    def __init__(self, filename):
        self.df = pd.read_csv(filename)
        self.df_normalized = None
        self.data_train = None
        self.data_test = None
        self.target_train = None
        self.target_test = None
        self.clf = None

    def preprocessing(self):
        # this can be done because the target value is either 0 or 1
        # thus apply the normalization have no effect on them
        scaler = MinMaxScaler()
        self.df_normalized = pd.DataFrame(scaler.fit_transform(self.df.values),
                                           columns=self.df.columns, index=self.df.index)
        self.split_dataframe()

    def split_dataframe(self):
        data = self.df_normalized[
            [col for col in self.df_normalized.columns if col != 'Survived']]
        target = self.df_normalized['Survived']
        # split the test and train data set
        self.data_train, self.data_test, self.target_train, self.target_test =
            train_test_split(data, target,
                             test_size=0.3,
                             shuffle=False)

    def fit_decision_tree(self):
        self.clf = tree.DecisionTreeClassifier()
        self.clf = self.clf.fit(self.data_train, self.target_train)

        # dot_data = tree.export_graphviz(clf, out_file=None)
        # graph = graphviz.Source(dot_data)
        # graph.format = 'jpg'
        # graph.render("decision_tree_plot_orignial")
        labels_test = self.clf.predict(self.data_test)
        acc = accuracy_score(labels_test, self.target_test)
        print("acc for test set is : " + str(acc))
        labels_test2 = self.clf.predict(self.data_train)
        acc2 = accuracy_score(labels_test2, self.target_train)
        print("acc for train set is : " + str(acc2))

    def find_optimal_decision_tree(self):
        # min_samples_leaf the minmum number of leaves a split
        # can happen according to the value of entropy
        auc_tain = {}
        auc_test = {}
        for i in range(2, 21):
            descion_tree = tree.DecisionTreeClassifier(min_samples_leaf=i)
            descion_tree.fit(self.data_train, self.target_train)
            false_positive_rate, true_positive_rate, thresholds = roc_curve(self.
                                     target_train, descion_tree.predict(self.data_train))

```

```

auc_tain[i] = auc(false_positive_rate, true_positive_rate)
false_positive_rate, true_positive_rate, thresholds = roc_curve(self.
    target_test, descion_tree.predict(self.data_test))
auc_test[i] = auc(false_positive_rate, true_positive_rate)
# fig = plt.figure()
# d = {'min_samples_leaf': np.array(list(auc_tain)), 'AUC score': np.array(
    list(auc_tain.values()))}
# pd_plot = pd.DataFrame(d)
# sns.lineplot(x='min_samples_leaf', y='AUC score', data=pd_plot)
# plt.show()
# fig.savefig('plot_train.png')
self.clf = tree.DecisionTreeClassifier(min_samples_leaf=6)
self.clf.fit(self.data_train, self.target_train)
# dot_data = tree.export_graphviz(self.clf, out_file=None)
# graph = graphviz.Source(dot_data)
# graph.format = 'jpg'
# graph.render("decision_tree_plot_optimal")

def part_D_calculation(self):
    print("female and survived: ")
    print(self.df.query('Sex == 1 & Survived ==1'))
    print("all female: ")
    print(self.df.query('Sex == 1'))
    print("first class and survived: ")
    print(self.df.query('Pclass == 1 & Survived ==1'))
    print("all first class: ")
    print(self.df.query('Pclass == 1'))

def print_df_normalized(self):
    print(self.df_normalized.head().to_string())

if __name__ == '__main__':
    titianic_sinking_model = TitanicSinkingModel("titanic.csv")
    titianic_sinking_model.preprocessing()
    # titianic_sinking_model.print_df_normalized()
    # print(titianic_sinking_model.data_train.to_string())
    titianic_sinking_model.fit_decision_tree()
    titianic_sinking_model.find_optimal_decision_tree()
    titianic_sinking_model.part_D_calculation()

```